

Similarity

Binary Similarity – Simple Matching Coefficient

- Simple Matching Coefficient (SMC)
- $SMC = \frac{M_{11} + M_{00}}{M_{01} + M_{10} + M_{11} + M_{00}}$
- M_{11} : The number of times that both vectors have a 1
- M_{01} : The number of times that vector A has a 0 when vector B has a 1
- M_{10} : The number of times that vector A has a 1 when vector B has a 0
- M_{00} : The number of times that both vectors have a 0
- Example:
 - $A = [0, 0, 1, 1, 0, 0, 1, 1]$
 - $B = [0, 1, 1, 0, 0, 1, 1, 1]$
 - $M_{11} = 3$
 - $M_{01} = 2$
 - $M_{10} = 1$
 - $M_{00} = 2$
 - $SMC = \frac{3+2}{2+1+3+2} = \frac{5}{8}$

Jaccard vs. SMC

- How do these two largely differ?
- $A = [0, 0, 0, 0, 0, 0, 0, 0, 0, 1]$
- $B = [1, 0, 0, 0, 0, 0, 0, 0, 0, 1]$
- $J = \frac{1}{1+0+1} = \frac{1}{2} = 0.5$
- $SMC = \frac{1+8}{1+0+1+8} = \frac{9}{10} = 0.9$
- Jaccard only considers the evidence of a 1 as important

Is 0 important?

- Scenario 1: You and I bump into each other at the grocery store and I look into both of our carts and I realized that we both did not buy the same stuff
 - My cart = $[1, 1, 1, 1, 0, 0, 0, 0, 0, 0, \dots, 0]$
 - Your cart = $[1, 1, 0, 0, 0, 0, 1, 0, 1, 0, \dots, 0]$
 - Where each index in the vectors represents a single item and 1 indicates that the item is in the cart and 0 represents the item is not in the cart
 - Would you say that our carts are similar?
 - * No


- Scenario 2: You and I are comparing our movies that we both have watched and state whether we like them (1) or dislike them (0)
 - My ratings = [1, 1, 0, 0, 0, 0, 0]
 - Your ratings = [0, 1, 0, 0, 0, 0, 0]
 - Would you say that we have similar preferences in movies?
 - * Yes

Asymmetric Attribute

- Binary attributes where the different values have different importance is considered an asymmetric attribute
- For example, where the presence of something (items in a grocery cart) gives us more information than what items are not in the grocery cart
- Jaccard captures this in a way by not considering the case with 0 in both vectors and only focusing on cases where at least one item / (1) is present

Values don't always play nice

- What data point is closest to B based on Euclidean distance?



- $d(B, A) = 100$
- $d(B, C) = 8$
- If one attribute has values that are magnitudes larger than another, it could be the controlling / dominating value in a distance calculation.
- This can also affect machine learning models

Scaling

- We normalize values to alleviate the issue of attributes with large values taking over
- Given attribute X , we can normalize our attribute using different techniques
 1. Min-Max Scaling / Normalization

$$\bullet X_{\text{new}} = \frac{X - \min(X)}{\max(X) - \min(X)}$$


2. Standardization (Z-Score Normalization)

- $X_{\text{new}} = \frac{X - \mu}{\sigma_X}$

Back to our example

- Assume that our attribute along the y -axis has values ranging from 50 to 1000
- Assume that our attribute along the x -axis has values ranging from 0 to 10
- Using the Min-Max Normalization

- $A_y = \frac{A_y - \min(y)}{\max(y) - \min(y)} = \frac{900 - 500}{1000 - 500} = \frac{400}{500} = 0.8$
- $A_x = \frac{A_x - \min(x)}{\max(x) - \min(y)} = \frac{1 - 0}{10 - 0} = \frac{1}{10} = 0.1$
- $B_y = \frac{B_y - \min(y)}{\max(y) - \min(y)} = \frac{800 - 500}{1000 - 500} = \frac{300}{500} = 0.6$
- $B_x = \frac{B_x - \min(x)}{\max(x) - \min(y)} = \frac{1 - 0}{10 - 0} = \frac{1}{10} = 0.1$
- $C_y = \frac{C_y - \min(y)}{\max(y) - \min(y)} = \frac{800 - 500}{1000 - 500} = \frac{300}{500} = 0.6$
- $C_x = \frac{C_x - \min(x)}{\max(x) - \min(y)} = \frac{9 - 0}{10 - 0} = \frac{9}{10} = 0.9$



- Now, what data point is closest to B based on Euclidean distance?
- $d(B, A) = 0.2$
- $d(B, C) = 0.7$
- A is now closer to B as we would think

K-Nearest Neighbours (KNN)

KNN Example for later


Instance / Object	a_1	a_2	y (label / class)
x_1	2	2	-1
x_2	1.5	1	-1
x_3	4.5	3	+1
x_4	4	4	-1
x_5	5	4.5	+1
x_6	3	6	+1

Eager vs. Lazy Classifiers


- Eager Learner:
 - Decision tree-based classifier
 - Eager learners: A model learns to map the input attributes to the class label
- Lazy Learner:
 - Does not build models explicitly
 - Memorizes entire training data and performs the majority of computations when a sample is given for classification
 - k -nearest neighbour classifier (KNN) is a lazy learner.

Nearest Neighbour Classifiers

- Basic Idea: If it walks like a duck, quacks like a duck, then it's probably a duck
- Requires 3 things
 - The training set
 - Distance metric to compute distance between the test object and each training object
 - The value of k , the number of nearest neighbours



- Given a training set (18 - and 13 +) and a test object (circle), there are three steps to classify the test object
 - Compute the distance between the test object and each training record
 - Identify the k nearest neighbours
 - Use class labels of the nearest neighbours to determine the class label of the test record (e.g., by taking majority vote)



- The k -nearest neighbours of a test instance x are the k training records that have the smallest distance to x

Choice of Similarity Measure Matters

- For example, for the word count of documents, cosine is better than correlation or Euclidean
- Consider the following pairs


$[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0]$ and $[0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]$

$[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1]$ and $[1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]$

- The Euclidean distance is 1.4142 for both pairs, but the cosine similarity measure has different values for these pairs

Nearest Neighbour Classification

- Data preprocessing is often required
 - Attributes may have to be scaled / normalized to prevent distance measures from being dominated by one of the attributes
 - * Height of a person may vary from 1.5m to 1.8m
 - * Weight of a person may vary from 90lb to 300lb
 - * Income of a person may vary from \$50K to \$1M
 - We can perform one of the previously discussed methods of normalization:
 - * Min-Max Normalization
 - * Standardization
- Choosing the value of k :
 - If k is too small, it is sensitive to noise points
 - If k is too large, the neighbourhood may include points from other classes
 - Typically, start with $k = \sqrt{n}$, and tune from there using dev. classes



- k -nearest neighbours can produce decision boundaries of arbitrary shapes
- Below is a Voronoi Diagram, with $k = 1$ for some dataset

